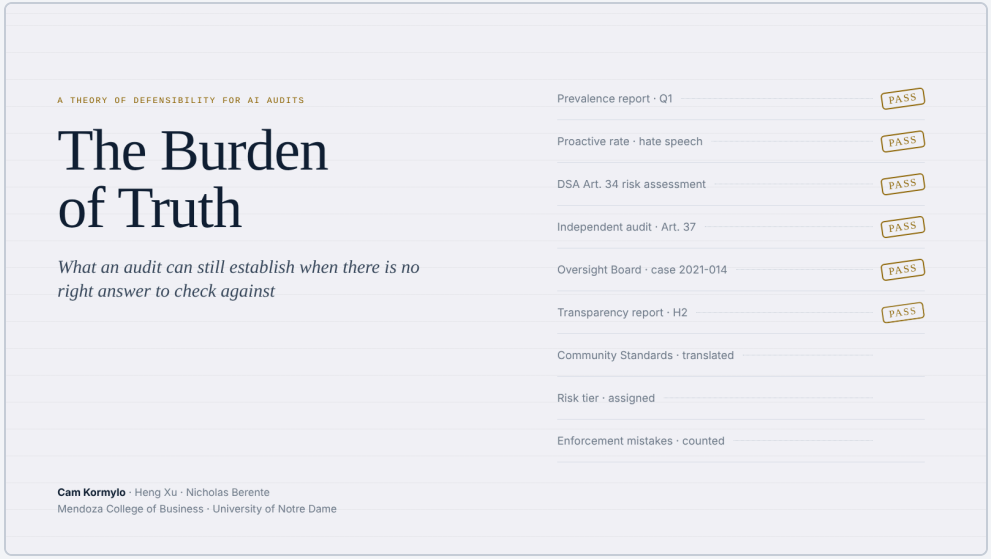
SPEAKER SCRIPT · 15-MINUTE TALK

The Burden of Truth: A Theory of Defensibility for AI Audits

Cam Kormylo, Heng Xu, Nicholas Berente · Mendoza College of Business, University of Notre Dame

Each page shows a slide in its final state on the left and the script for that slide on the right. The small green CLICK marks show where to advance the slide. Time cues are cumulative and assume a pace of about 155 words a minute; the full script is about 2511 words, which runs about 16 minutes read verbatim and a little less spoken from memory. If running long, the slides to shorten are 2, 13 and 14.

Slides: ckormylo.github.io/burden-of-truth



SLIDE 1 OF 16

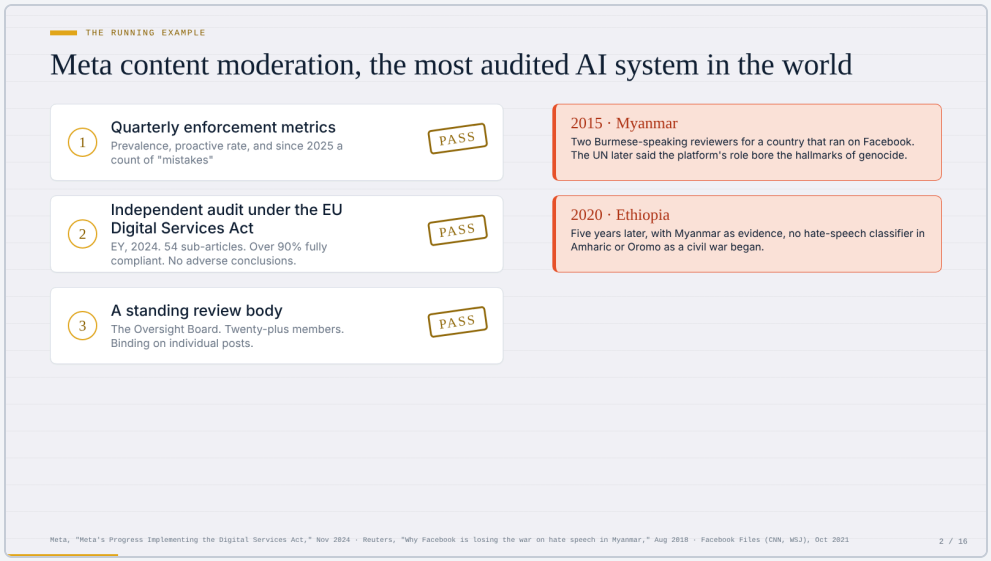
0:00 · ABOUT 41 SECONDS

The Burden of Truth

Thanks for having me. This talk is about auditing AI systems, and specifically about what an audit can still tell us when there is no right answer to check the system against. It comes from a paper with Heng Xu and Nick Berente at Notre Dame.

I will use one running example throughout, Meta's content moderation. The argument is simple to state. The way we audit AI is inherited from financial auditing. That inheritance assumes we can check a system's outputs against a standard. For a growing class of AI systems, no such standard exists. So we propose a different question for the audit to ask.

106 words



SLIDE 2 OF 16

0:41 · ABOUT 95 SECONDS · 5 CLICKS

Meta content moderation, the most audited AI system in the world

Here is the example. Meta decides which posts stay up and which come down on Facebook and Instagram, for billions of users, in well over a hundred languages. Most of that work is done by automated systems, and as far as I can tell it is the most audited AI system in operation.

CLICK Every quarter Meta publishes its own enforcement numbers: how much violating content is on the platform, what share of removals its systems caught on their own, and since last year, how many mistakes it made.

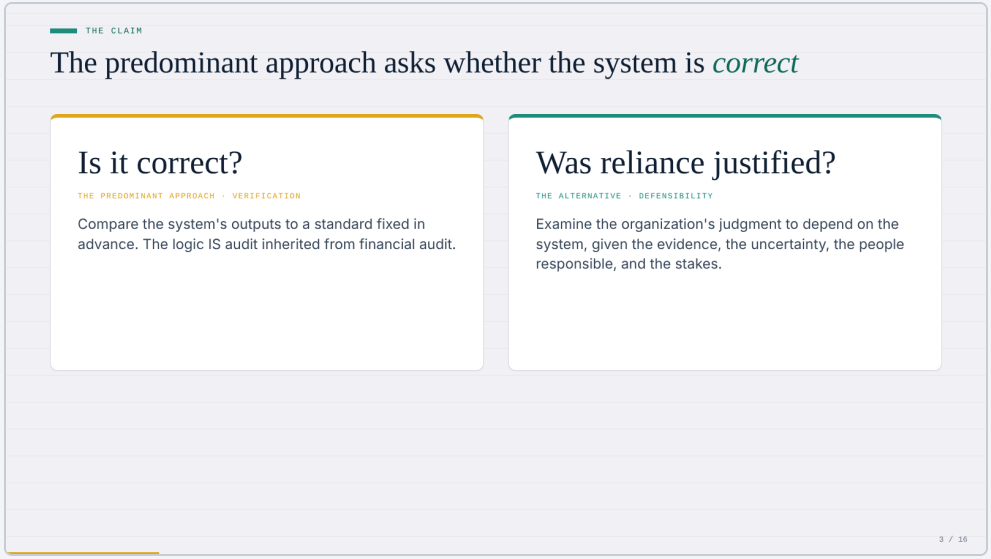
CLICK Under the European Digital Services Act, an outside firm, EY, audited Meta's compliance in 2024 across 54 sub-articles of the law and found over ninety percent fully compliant, with no adverse conclusions.

CLICK And Meta has an Oversight Board, an independent body that can reverse Meta's decisions on individual posts. That is a serious audit regime, and I am not here to mock it.

CLICK And yet. In 2015, Facebook had two people reviewing content in Burmese, for a country where Facebook effectively was the internet. The UN later found that the platform played a determining role in the violence against the Rohingya.

CLICK Five years later, in Ethiopia, as a civil war began, Meta had no hate-speech classifier for Amharic or Oromo, the country's two largest languages, with Myanmar already on the record. Every audit in that stack passed, before and after both of these. I want to talk about what they were checking.

245 words



SLIDE 3 OF 16

2:16 · ABOUT 65 SECONDS · 1 CLICK

The predominant approach asks whether the system is correct

The predominant approach to auditing a system is to ask whether it is correct. You compare what the system produces to a standard that was fixed in advance. That is verification, and it is the logic that information systems auditing inherited from financial auditing. The ledger either reconciles or it does not. That is the question the Meta audits I just showed you were asking, and it is the question nearly every AI audit framework asks.

CLICK We propose an alternative. Ask instead whether the organization was justified in relying on this system. That looks at the organization's judgment to depend on the system, given the evidence it had, the uncertainty that remained, the people responsible, and the stakes of being wrong. We call that defensibility.

The argument of the paper is that for a large class of AI systems the predominant question cannot be answered and the alternative can. First why the predominant approach stops working, then what the alternative requires, then both audits run on one deployment.

169 words



SLIDE 4 OF 16

3:21 · ABOUT 47 SECONDS · 4 CLICKS

Verification rests on four requirements

Start with why verification works where it works. The 1972 definition of auditing, which the profession still uses, presupposes four things.

CLICK Assertions that can be true or false. "The balance is ten million."

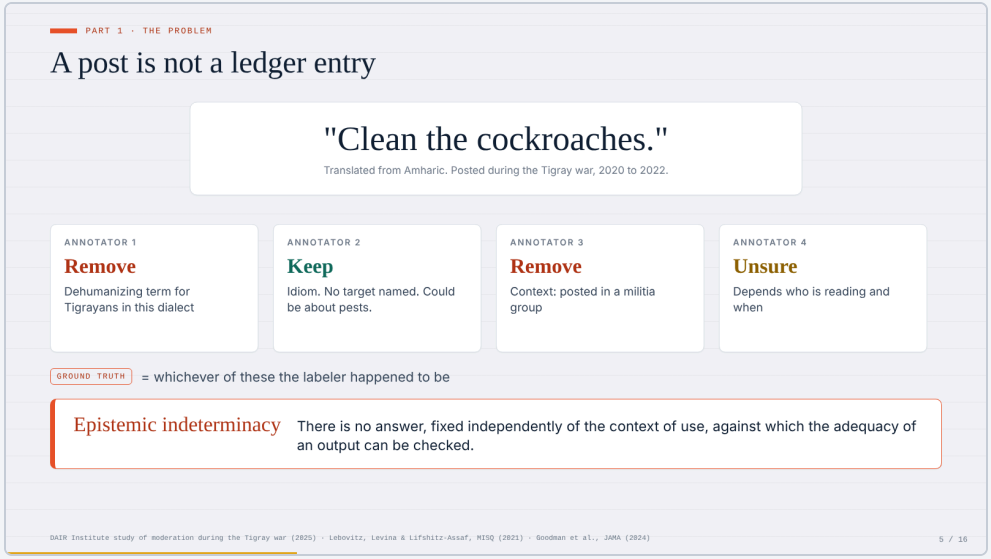
CLICK Evidence that can confirm or disconfirm them.

CLICK A subject matter with determinate states, meaning there is a fact of the matter to compare against.

CLICK And criteria set before the audit begins.

When all four hold, an audit can detect and report nonconformance, and some content moderation is like this. A known child-abuse image either matches the hash database or it does not. Every AI audit framework of the last decade, up to the EU AI Act, widened what gets checked without changing this logic.

121 words



SLIDE 5 OF 16

4:08 · ABOUT 81 SECONDS · 3 CLICKS

A post is not a ledger entry

Now a post. CLICK This one is from the Tigray war in Ethiopia, translated from Amharic by a research group that studied moderation during that conflict: "Clean the cockroaches."

CLICK They gave posts like this to fluent annotators, people who speak the language and know the context, and the annotators disagreed. One reads it as a dehumanizing term for Tigrayans in that dialect. One reads it as an idiom with no target. One says it depends on where it was posted. One says it depends on who is reading and when. There is no fact of the matter for a classifier to be checked against. What we call the ground truth is whichever of these people happened to do the labeling.

CLICK This is the condition the paper is about. We call it epistemic indeterminacy: there is no answer, fixed independently of the context of use, against which the adequacy of an output can be checked. It removes the third requirement, determinate states, and the other three fall with it. This is not a defect that a better model fixes. It is a property of the task, and it shows up anywhere an output's adequacy depends on who uses it: a hospital discharge summary, a translation, a legal summary.

209 words



SLIDE 6 OF 16

5:29 · ABOUT 59 SECONDS · 4 CLICKS

Two common responses

There are two standard responses to this.

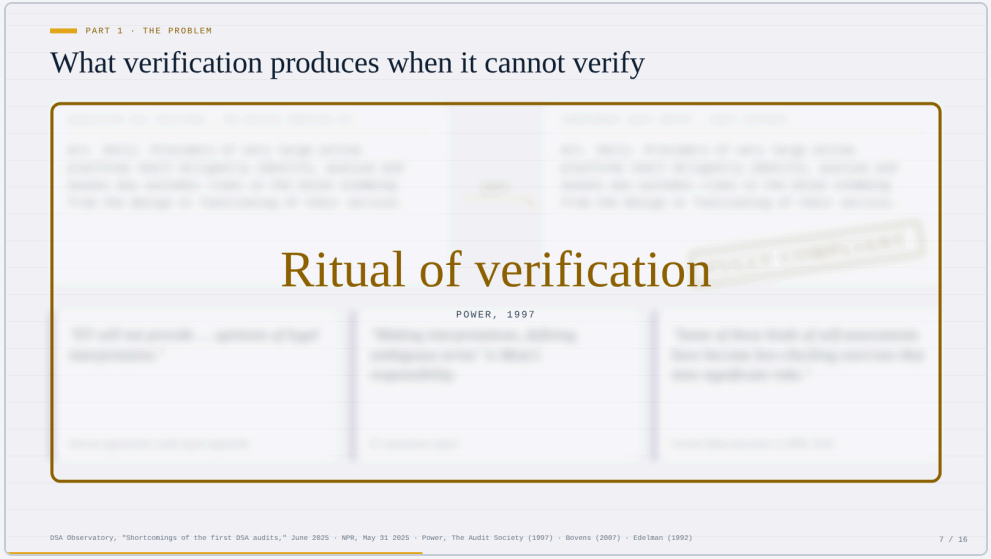
CLICK The first is to treat it as a measurement problem. Meta told the Financial Times this year that its language models make thirteen percent fewer errors than human reviewers.

CLICK Fewer errors measured against what? Against human labels. The yardstick is made of the thing in dispute.

CLICK The second response is to wait, because the systems are improving. In the first quarter of 2025, 97.4 percent of hate-speech removals on Instagram were automated.

CLICK But improvement makes verification harder, not easier. To show that an error happens less than once in a thousand cases, you need about three thousand clean cases in a row. And the rarer the error, the less attention the remaining human reviewers pay to a system that is usually right. So verification costs more and catches less at exactly the point where the organization relies on the system more.

152 words



SLIDE 7 OF 16

6:28 · ABOUT 61 SECONDS · 4 CLICKS

What verification produces when it cannot verify

What does verification produce when it cannot verify? Here is the DSA audit I mentioned, as a group of legal scholars at the DSA Observatory read it.

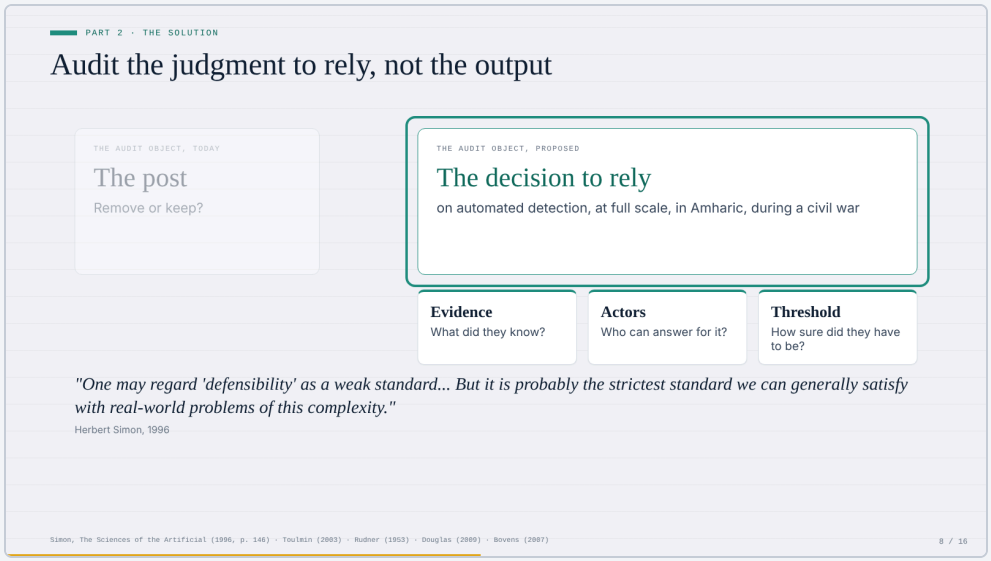
CLICK The audit criteria are the text of the statute, copied.

CLICK The conclusion is "fully compliant."

CLICK And who decides what the law's ambiguous terms mean? The audit agreement says EY will not provide opinions of legal interpretation, and that defining ambiguous terms is Meta's responsibility. So the auditee writes the criteria and the auditor confirms correspondence. The third quote, from a former Meta executive this year, says the same about a different internal review.

CLICK Michael Power has a name for this: a ritual of verification. The audit continues, but it stops informing anyone. And no bad faith is required. A reviewer who cannot judge whether the system is reliable judges what it can see, which is whether the paperwork is complete, and organizations respond rationally by producing paperwork.

158 words



SLIDE 8 OF 16

7:29 · ABOUT 77 SECONDS · 3 CLICKS

Audit the judgment to rely, not the output

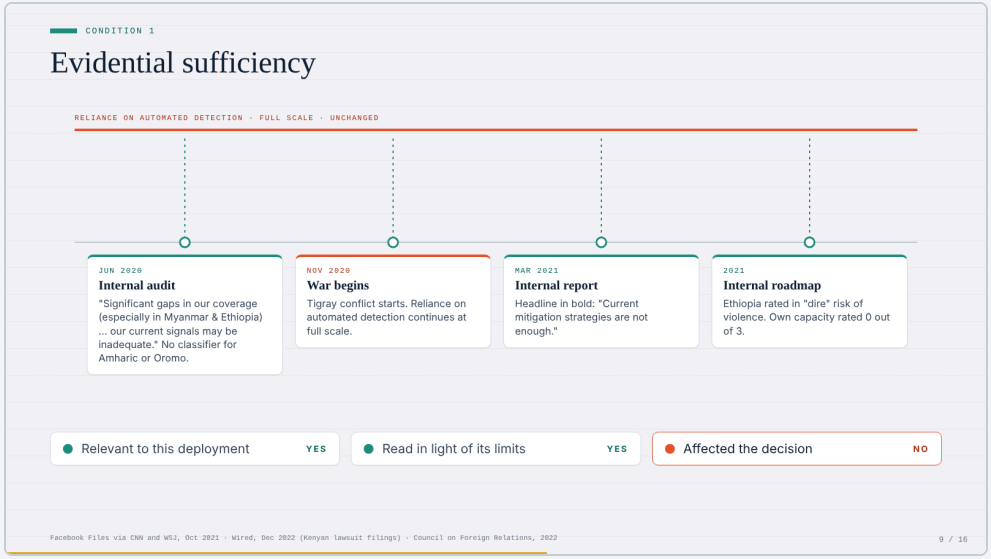
So here is the move. Every check rests on prior choices: what to check, against what criteria, and how much checking is enough. For a hash database those choices are a specification, made once. For Amharic during a civil war, they are judgments, made by particular people about a particular deployment.

CLICK That is where the organization's responsibility actually lives, and it is something an auditor can examine even when correctness is unavailable, because it concerns the quality of a judgment rather than a comparison between an output and a fact. So we shift the object of the audit from the output to the decision to rely on the system. Verification stays, as one input to that decision.

CLICK A justification for relying on a system has three parts: the evidence it rests on, the actors who stand behind it, and the level of confidence it claims the evidence meets. Each replaces one of the requirements of verification, and each becomes a condition of defensibility.

CLICK Herbert Simon anticipated the objection. One may regard defensibility as a weak standard, he wrote, but it is probably the strictest standard we can generally satisfy with real-world problems of this complexity.

198 words



SLIDE 9 OF 16

8:46 · ABOUT 65 SECONDS · 5 CLICKS

Condition 1 · Evidential sufficiency

Condition one is evidential sufficiency. The question is whether the evidence bore on the deployment actually being made. Was it relevant to this deployment, was it read with its limits in mind, and did it affect the decision. The flat line across the top is Meta's reliance on automated detection in Ethiopia, at full scale. It never moves.

CLICK In June 2020, Meta's own internal audit found significant gaps in coverage, especially in Myanmar and Ethiopia, and no classifier for Amharic or Oromo.

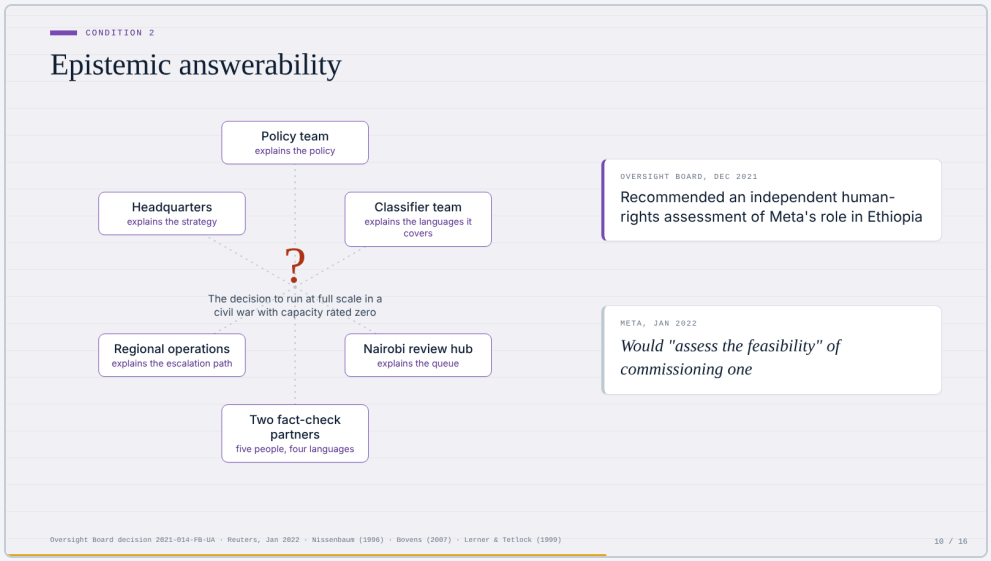
CLICK In November 2020, the war began.

CLICK In March 2021, an internal report put it in a bold headline: current mitigation strategies are not enough.

CLICK And the 2021 roadmap rated Ethiopia at dire risk of violence and rated Meta's own capacity there zero out of three.

CLICK So the evidence was relevant, and it was read. It did not change the decision. Two of the three clauses met. Verification never tests the third, because the third is about the judgment, not the data.

168 words



SLIDE 10 OF 16

9:51 · ABOUT 60 SECONDS · 4 CLICKS

Condition 2 · Epistemic answerability

Condition two is epistemic answerability. The question is whether identifiable people, who could have shaped or halted the deployment, can defend the judgment as a whole, including what would have changed it.

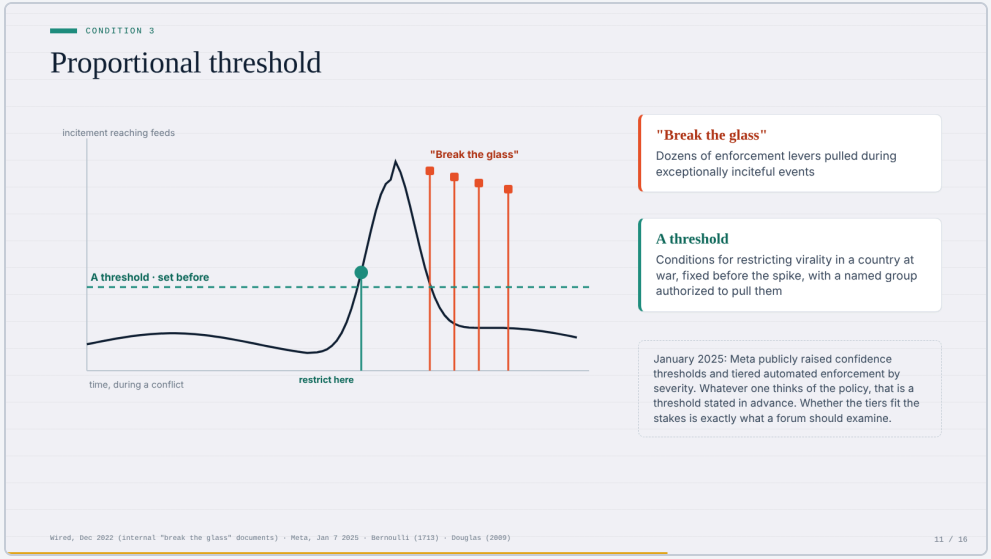
CLICK At Meta, everyone can explain their own part. The policy team explains the policy. The classifier team explains which languages it covers. The review hub in Nairobi explains the queue. Two fact-checking partners, five people between them, cover four languages.

CLICK Nobody is positioned to answer for the center: the decision to operate at full scale in a civil war with capacity rated zero. Helen Nissenbaum called this the problem of many hands.

CLICK A forum did ask. In December 2021 the Oversight Board recommended an independent human-rights assessment of Meta's role in Ethiopia.

CLICK Meta said it would assess the feasibility of commissioning one. That is the characteristic failure of this condition: nominal responsibility without the capacity to give an account.

155 words



SLIDE 11 OF 16

10:51 · ABOUT 67 SECONDS · 4 CLICKS

Condition 3 · Proportional threshold

Condition three is proportional threshold. The question is whether the confidence required before relying on the system was fixed in advance, scaled to how bad and how reversible the harm could be, and backed by something that could actually restrict the system.

CLICK Incitement during a conflict does not arrive on a schedule. It spikes.

CLICK Meta's instruments for this are what its internal documents call break-the-glass measures: dozens of enforcement levers pulled during exceptionally inciteful events. They are pulled after the spike. That is a threshold inferred from results.

CLICK A proportional threshold is drawn beforehand. Here is the level, here is the named group that pulls the lever, here is what the lever does.

CLICK One contrast, to be fair. In January 2025 Meta publicly raised the confidence its systems need before removing content and tiered automated enforcement by severity. Whatever you think of that policy, it is a threshold stated in advance. Whether the tiers fit the stakes is exactly what a forum should examine, and nobody is positioned to.

173 words



SLIDE 12 OF 16

11:58 · ABOUT 38 SECONDS · 3 CLICKS

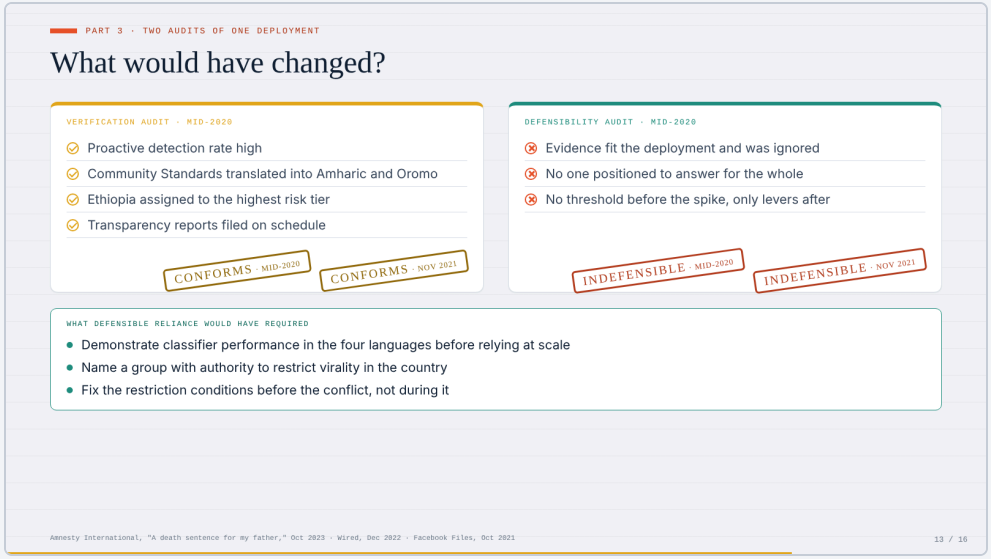
Three parts of one justification

CLICK The three conditions are parts of a single justification.

CLICK Remove any one and it does not stand. Strong evidence does not make up for nobody who can answer for it. A careful threshold does not make up for a process that never engaged the evidence.

CLICK A justification missing a part is incomplete, which is a different defect from being weak. This is not a scorecard that sums across dimensions. And a finding of defensible reliance does not mean the platform is safe. It means the decision was justified on what was knowable when it was made.

99 words



SLIDE 13 OF 16

12:36 · ABOUT 93 SECONDS · 5 CLICKS

What would have changed?

Now put the two audits side by side, on the same deployment, in the middle of 2020.

CLICK A verification audit finds everything it can check in order. It conforms.

CLICK A defensibility audit finds the justification fails on all three conditions, at a time when every reported metric looked fine. Indefensible.

CLICK Now move forward to November 2021. A chemistry professor named Meareg Amare is shot outside his home after posts naming him and giving his address circulated on Facebook for weeks. Run both audits again.

CLICK Verification passes again, because nothing it checks has changed. Defensibility returns the same finding it gave before, because it did not need the harm to identify the reliance that produced it. Neither finding depends on how things turned out.

CLICK And the defensibility finding says what would have changed. Demonstrate that the classifiers work in the four languages before relying on them at scale. Name a group with the authority to restrict virality in the country. Fix the conditions for doing that before the conflict rather than during it. Those are the measures Meta eventually took, and the ones Amnesty says came too late. I am not claiming the war would have gone differently. I am claiming the reliance that produced this failure was identifiable before the failure, and that the fix was known. The forum that did eventually form was a court in Kenya, in a suit brought by the professor's son.

241 words



SLIDE 14 OF 16

14:09 · ABOUT 54 SECONDS · 3 CLICKS

The judgment is being made again

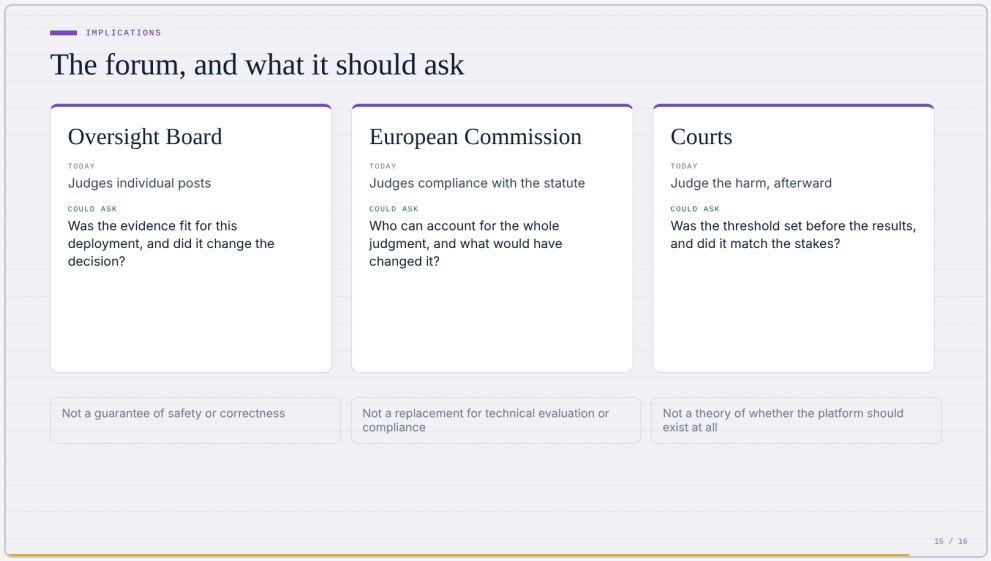
This is not history.

CLICK In 2025 Meta replaced roughly half of its human content review requests with large language models, and its target is over ninety percent for some content types by the end of this year. Staff told the Financial Times the rollout is too fast to oversee.

CLICK And the risk review itself is being automated. According to internal documents reported by NPR, up to ninety percent of Meta's product risk assessments will be done by an AI system. Product teams fill in a questionnaire and receive an instant decision, including for AI safety, youth risk, and integrity.

CLICK The paper's second proposition says that as reliability rises, the evidence needed to verify grows and human scrutiny falls, so the justification for reliance needs more support, not less. Reliance is renewed every day. Is the justification?

139 words



SLIDE 15 OF 16

15:03 · ABOUT 50 SECONDS · 3 CLICKS

The forum, and what it should ask

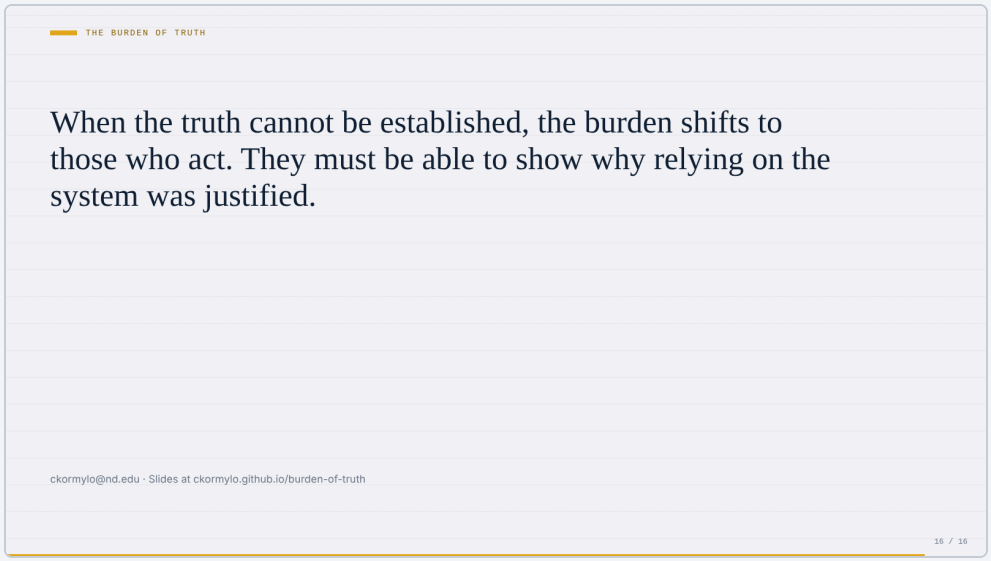
Three forums already sit in judgment of Meta.

CLICK The Oversight Board judges individual posts. The European Commission judges compliance with the statute. Courts judge the harm, afterward.

CLICK Each could ask one question per condition. Was the evidence fit for this deployment, and did it change the decision? Who can account for the whole judgment, and what would have changed it? Was the threshold set before the results, and did it match the stakes? None of them is positioned to ask those today.

CLICK Three limits, briefly. Defensibility is not a guarantee of safety or correctness. It does not replace technical evaluation or legal compliance. And it is not a theory of whether a platform should exist at all. That is a democratic question the audit does not decide.

130 words



SLIDE 16 OF 16

15:53 · ABOUT 19 SECONDS

Close

When the truth cannot be established, the burden shifts to those who act. They must be able to show why relying on the system was justified.

Meta passed every audit it was given. Nobody gave it the one that asked whether anyone was right to rely.

Thank you.

48 words